

model

February 28, 2026

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
import shap
from sklearn.utils import shuffle
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.calibration import CalibratedClassifierCV, calibration_curve
from sklearn.metrics import (
    roc_auc_score, accuracy_score, roc_curve, brier_score_loss,
    precision_score, recall_score, f1_score, confusion_matrix,
    ↪classification_report
)

# Load and prepare training data
train_data = pd.read_csv("training_dataset.csv")
X_train = train_data.drop(columns=["Diagnosed", "Anxiety Level (1-10)"])
y_train = train_data["Diagnosed"]
X_train, y_train = shuffle(X_train, y_train, random_state=42)

# Load and prepare test data
test_data = pd.read_csv("testing_dataset.csv")
X_test = test_data.drop(columns=["Diagnosed", "Anxiety Level (1-10)"])
y_test = test_data["Diagnosed"]

# Identify feature types
numeric_features = X_train.select_dtypes(include=[float, int]).columns.tolist()
categorical_features = X_train.select_dtypes(include=["object", "category"]).
    ↪columns.tolist()

# Create preprocessing pipeline
```

```

preprocessor = ColumnTransformer([
    ("num", StandardScaler(), numeric_features),
    ("cat", OneHotEncoder(drop="first", sparse_output=False),
    ↪categorical_features)
])

# Create model pipelines
logreg_pipeline = Pipeline([
    ("preprocessor", preprocessor),
    ("classifier", LogisticRegression(max_iter=1000, class_weight="balanced",
    ↪solver="saga", l1_ratio=0.5, C=1.0))
])

rf_pipeline = Pipeline([
    ("preprocessor", preprocessor),
    ("classifier", RandomForestClassifier(n_estimators=200,
    ↪class_weight="balanced", random_state=42))
])

scale_pos_weight = (y_train == 0).sum() / (y_train == 1).sum()
xgb_pipeline = Pipeline([
    ("preprocessor", preprocessor),
    ("classifier", XGBClassifier(n_estimators=200, max_depth=4, learning_rate=0.
    ↪1,
                                eval_metric="logloss",
    ↪scale_pos_weight=scale_pos_weight, random_state=42))
])

# Calibrate models
logreg_model = CalibratedClassifierCV(logreg_pipeline, method="isotonic", cv=5)
rf_model = CalibratedClassifierCV(rf_pipeline, method="isotonic", cv=5)
xgb_model = CalibratedClassifierCV(xgb_pipeline, method="isotonic", cv=5)

# Train models
logreg_model.fit(X_train, y_train)
rf_model.fit(X_train, y_train)
xgb_model.fit(X_train, y_train)

# Get predicted probabilities
logreg_prob = logreg_model.predict_proba(X_test)[: , 1]
rf_prob = rf_model.predict_proba(X_test)[: , 1]
xgb_prob = xgb_model.predict_proba(X_test)[: , 1]

# Show first 20 predictions
print("\nPredicted percentages (first 20 patients):")
df_probs = pd.DataFrame({
    "Patient": range(1, 21),

```

```

        "Diagnosis(1-Anxiety, 0-No Anxiety)": y_test.iloc[:20].values,
        "Logistic Regression (%)": np.round(logreg_prob[:20]*100, 2),
        "Random Forest (%)": np.round(rf_prob[:20]*100, 2),
        "XGBoost (%)": np.round(xgb_prob[:20]*100, 2)
    })
print(df_probs.to_string(index=False))

# Find optimal thresholds
def find_best_threshold(y_true, y_prob):
    fpr, tpr, thresholds = roc_curve(y_true, y_prob)
    best_idx = np.argmax(tpr - fpr)
    return thresholds[best_idx]

logreg_threshold = find_best_threshold(y_test, logreg_prob)
rf_threshold = find_best_threshold(y_test, rf_prob)
xgb_threshold = find_best_threshold(y_test, xgb_prob)

# Plot probability distributions
def plot_probabilities_horizontal(y_true, y_prob, threshold, name):
    df = pd.DataFrame({"Probability": y_prob, "Diagnosis": y_true})
    plt.figure(figsize=(10, 4))
    sns.histplot(data=df, x="Probability", hue="Diagnosis", bins=20,
        palette=["skyblue", "salmon"], alpha=0.6)
    plt.axvline(threshold, color="green", linestyle="--", label=f"Threshold_{threshold:.2f}")
    plt.xlabel("Predicted Probability")
    plt.ylabel("Count")
    plt.title(name)
    plt.legend()
    plt.show()

plot_probabilities_horizontal(y_test, logreg_prob, logreg_threshold, "Logistic_Regression")
plot_probabilities_horizontal(y_test, rf_prob, rf_threshold, "Random Forest")
plot_probabilities_horizontal(y_test, xgb_prob, xgb_threshold, "XGBoost")

# feature importance function
def plot_feature_importance(pipeline_model, name):

    classifier = pipeline_model.named_steps["classifier"]

    preprocessor_temp = ColumnTransformer([
        ("num", StandardScaler(), numeric_features),
        ("cat", OneHotEncoder(drop="first", sparse_output=False),
        categorical_features)
    ])
    preprocessor_temp.fit(X_train)

```

```

numeric_feature_names = numeric_features
cat_feature_names = preprocessor_temp.named_transformers_["cat"].
↳get_feature_names_out(categorical_features)
feature_names = list(numeric_feature_names) + list(cat_feature_names)

plt.figure(figsize=(12, 8))

if hasattr(classifier, "feature_importances_"):
    importances = classifier.feature_importances_
    indices = np.argsort(importances)[-15:]
    plt.barh(range(len(indices)), importances[indices])
    plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
    plt.xlabel("Feature Importance")
    plt.title(f"{name} - Feature Importance (Top 15)")
    print(f"\nTop 15 features for {name}:")
    for idx in reversed(indices):
        print(f" {feature_names[idx]}: {importances[idx]:.4f}")

elif hasattr(classifier, "coef_"):
    coefficients = classifier.coef_[0]
    abs_coef = np.abs(coefficients)
    indices = np.argsort(abs_coef)[-15:]
    colors = ['red' if c < 0 else 'green' for c in coefficients[indices]]
    plt.barh(range(len(indices)), coefficients[indices], color=colors)
    plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
    plt.xlabel("Coefficient Value")
    plt.title(f"{name} - Feature Coefficients (Top 15)")
    from matplotlib.patches import Patch
    legend_elements = [Patch(facecolor='green', label='Positive (increases_
↳risk)'),
                        Patch(facecolor='red', label='Negative (decreases_
↳risk)')]
    plt.legend(handles=legend_elements, loc='lower right')
    print(f"\nTop 15 features for {name}:")
    for idx in reversed(indices):
        direction = "increases risk" if coefficients[idx] > 0 else_
↳"decreases risk"
        print(f" {feature_names[idx]}: {coefficients[idx]:.4f}_
↳({direction})")

plt.axvline(x=0, color='black', linestyle='-', linewidth=0.5)
plt.tight_layout()
plt.show()

logreg_pipeline.fit(X_train, y_train)

```

```

rf_pipeline.fit(X_train, y_train)
xgb_pipeline.fit(X_train, y_train)

# Plot feature importance for all models
plot_feature_importance(logreg_pipeline, "Logistic Regression")
plot_feature_importance(rf_pipeline, "Random Forest")
plot_feature_importance(xgb_pipeline, "XGBoost")

# evaluation model
def evaluate_model(y_true, y_prob, threshold, name):
    y_pred = (y_prob >= threshold).astype(int)
    tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    roc_auc = roc_auc_score(y_true, y_prob)
    brier = brier_score_loss(y_true, y_prob)
    sensitivity = tp / (tp + fn) if (tp + fn) > 0 else 0
    specificity = tn / (tn + fp) if (tn + fp) > 0 else 0

    print(f"{name} METRICS (threshold={threshold:.3f})")
    print(f"Confusion Matrix:")
    print(f"  True Negatives: {tn}")
    print(f"  False Positives: {fp}")
    print(f"  False Negatives: {fn}")
    print(f"  True Positives: {tp}")
    print(f"\nKey Performance Metrics:")
    print(f"  • SENSITIVITY (Recall/True Positive Rate): {sensitivity:.4f}␣
↪({sensitivity*100:.2f}%)")
    print(f"  • SPECIFICITY (True Negative Rate): {specificity:.4f}␣
↪({specificity*100:.2f}%)")
    print(f"  • F1-SCORE: {f1:.4f}")
    print(f"\nAdditional Metrics:")
    print(f"  Accuracy: {acc:.4f} ({acc*100:.2f}%)")
    print(f"  Precision: {prec:.4f} ({prec*100:.2f}%)")
    print(f"  ROC AUC: {roc_auc:.4f}")
    print(f"  Brier Score: {brier:.4f}")
    print(f"\nClassification Report:")
    print(classification_report(y_true, y_pred, digits=4))

    return confusion_matrix(y_true, y_pred), {
        'sensitivity': sensitivity,
        'specificity': specificity,
        'f1_score': f1,
        'precision': prec,

```

```

        'accuracy': acc,
        'roc_auc': roc_auc
    }

# Evaluate models
logreg_cm, logreg_metrics = evaluate_model(y_test, logreg_prob,
    ↳logreg_threshold, "Logistic Regression")
rf_cm, rf_metrics = evaluate_model(y_test, rf_prob, rf_threshold, "Random
    ↳Forest")
xgb_cm, xgb_metrics = evaluate_model(y_test, xgb_prob, xgb_threshold, "XGBoost")

# Model comparison table
comparison_df = pd.DataFrame({
    'Model': ['Logistic Regression', 'Random Forest', 'XGBoost'],
    'Threshold': [logreg_threshold, rf_threshold, xgb_threshold],
    'Sensitivity (%)': [logreg_metrics['sensitivity']*100,
    ↳rf_metrics['sensitivity']*100, xgb_metrics['sensitivity']*100],
    'Specificity (%)': [logreg_metrics['specificity']*100,
    ↳rf_metrics['specificity']*100, xgb_metrics['specificity']*100],
    'F1-Score': [logreg_metrics['f1_score'], rf_metrics['f1_score'],
    ↳xgb_metrics['f1_score']],
    'Precision (%)': [logreg_metrics['precision']*100,
    ↳rf_metrics['precision']*100, xgb_metrics['precision']*100],
    'Accuracy (%)': [logreg_metrics['accuracy']*100,
    ↳rf_metrics['accuracy']*100, xgb_metrics['accuracy']*100],
    'ROC-AUC': [logreg_metrics['roc_auc'], rf_metrics['roc_auc'],
    ↳xgb_metrics['roc_auc']]
}).round(2)

print("\nMODEL COMPARISON SUMMARY\n", comparison_df.to_string(index=False))

# ROC curve
plt.figure(figsize=(8,6))
for prob, label in zip([logreg_prob, rf_prob, xgb_prob], ["Logistic
    ↳Regression", "Random Forest", "XGBoost"]):
    fpr, tpr, _ = roc_curve(y_test, prob)
    auc = roc_auc_score(y_test, prob)
    plt.plot(fpr, tpr, label=f"{label} (AUC={auc:.3f})")
plt.plot([0,1], [0,1], color='gray', linestyle='--')
plt.xlabel("False Positive Rate (1 - Specificity)")
plt.ylabel("True Positive Rate (Sensitivity)")
plt.title("ROC-AUC Curve Comparison")
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

```

```

# Confusion matrices
fig, axes = plt.subplots(1, 3, figsize=(15, 4))
sensitivity_labels = [
    f"Sensitivity: {logreg_metrics['sensitivity']:.2%}\nSpecificity:␣
    ↳{logreg_metrics['specificity']:.2%}",
    f"Sensitivity: {rf_metrics['sensitivity']:.2%}\nSpecificity:␣
    ↳{rf_metrics['specificity']:.2%}",
    f"Sensitivity: {xgb_metrics['sensitivity']:.2%}\nSpecificity:␣
    ↳{xgb_metrics['specificity']:.2%}"
]
for idx, (cm, ax, label) in enumerate(zip([logreg_cm, rf_cm, xgb_cm], axes,␣
    ↳sensitivity_labels)):
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax,
                xticklabels=['Predicted Negative', 'Predicted Positive'],
                yticklabels=['Actual Negative', 'Actual Positive'])
    ax.set_xlabel("Predicted")
    ax.set_ylabel("Actual")
    ax.set_title(f"[ 'Logistic Regression', 'Random Forest',␣
    ↳'XGBoost' ][idx]]\n{label}")
plt.tight_layout()
plt.show()

# Calibration curve
plt.figure(figsize=(15,6))
for prob, label in zip([logreg_prob, rf_prob, xgb_prob], ["Logistic␣
    ↳Regression", "Random Forest", "XGBoost"]):
    prob_true, prob_pred = calibration_curve(y_test, prob, n_bins=10)
    plt.plot(prob_pred, prob_true, marker='o', label=label)
plt.plot([0,1], [0,1], linestyle='--', color='gray')
plt.xlabel("Predicted Probability")
plt.ylabel("Observed Probability")
plt.title("Combined Calibration Curve")
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

# Get feature names after preprocessing
preprocessor_temp = ColumnTransformer([
    ("num", StandardScaler(), numeric_features),
    ("cat", OneHotEncoder(drop="first", sparse_output=False),␣
    ↳categorical_features)
])
preprocessor_temp.fit(X_train)

numeric_feature_names = numeric_features

```

```

cat_feature_names = preprocessor_temp.named_transformers_["cat"].
    ↪get_feature_names_out(categorical_features)
feature_names = list(numeric_feature_names) + list(cat_feature_names)

# Transform the data
X_train_transformed = preprocessor_temp.transform(X_train)
X_test_transformed = preprocessor_temp.transform(X_test)

# Get the underlying XGBoost model (not the calibrated one)
xgb_model_uncalibrated = xgb_pipeline.named_steps['classifier']

# Create SHAP explainer
explainer = shap.TreeExplainer(xgb_model_uncalibrated)

# Calculate SHAP values for test set
shap_values = explainer.shap_values(X_test_transformed)

# Create waterfall plots for only 3 patients (first 3)
print("\nGenerating SHAP Waterfall Plots for First 3 Patients...")
print("="*60)

for i in range(min(3, len(X_test))):
    plt.figure(figsize=(12, 7))

    # Get prediction and true value
    pred_prob = xgb_prob[i]
    true_diag = y_test.iloc[i]

    # Create waterfall plot
    shap.waterfall_plot(
        shap.Explanation(
            values=shap_values[i],
            base_values=explainer.expected_value,
            data=X_test_transformed[i],
            feature_names=feature_names
        ),
        show=False,
        max_display=12 # Show top 12 features for readability
    )

    # Add title with patient information
    plt.title(f"Patient {i+1} - {'DIAGNOSED' if true_diag == 1 else 'NOT_
    ↪DIAGNOSED'}\n"
            f"Predicted Probability: {pred_prob:.2%} | Base Value: {explainer.
    ↪expected_value:.3f}",
            fontsize=14, fontweight='bold')

```



```

plt.tight_layout()
plt.show()
print("-" * 80)

print("RESULTS TABLE")

# Function to get top features for a patient based on SHAP values
def get_top_shap_features(shap_values_row, feature_names, n_top=5,
    positive=True):
    """
    Get top n features with highest positive or negative SHAP values
    """
    if positive:
        # Get indices of top positive SHAP values
        indices = np.argsort(shap_values_row)[-n_top:][::-1]
    else:
        # Get indices of top negative SHAP values (most negative)
        indices = np.argsort(shap_values_row)[:n_top]

    features = []
    for idx in indices:
        if positive and shap_values_row[idx] > 0:
            features.append(f"{feature_names[idx]} (+{shap_values_row[idx]:.
    ↪3f})")
        elif not positive and shap_values_row[idx] < 0:
            features.append(f"{feature_names[idx]} ({shap_values_row[idx]:.
    ↪3f})")

    # Pad with empty strings if less than n_top
    while len(features) < n_top:
        features.append("")

    return features

# Create risk categories based on predicted probability
def categorize_risk(prob, low_thresh=0.3, high_thresh=0.7):
    if prob < low_thresh:
        return "Low Risk"
    elif prob < high_thresh:
        return "Moderate Risk"
    else:
        return "High Risk"

xgb_risk_category = [categorize_risk(p) for p in xgb_prob]

```

```

# Get top features for each patient
top_increase_features = []
top_decrease_features = []

for i in range(len(xgb_prob)):
    increase_features = get_top_shap_features(shap_values[i], feature_names,
    ↪n_top=5, positive=True)
    decrease_features = get_top_shap_features(shap_values[i], feature_names,
    ↪n_top=5, positive=False)

    top_increase_features.append(increase_features)
    top_decrease_features.append(decrease_features)

# Create enhanced results table
results_table = pd.DataFrame({
    "Patient": range(1, len(xgb_prob) + 1),
    "Diagnosis(1-Anxiety, 0-No Anxiety)": ["Diagnosed" if val == 1 else "Not",
    ↪Diagnosed" for val in y_test.values],
    "Predicted Prob (%)": np.round(xgb_prob * 100, 2),
    "Risk Category": xgb_risk_category,
    "Top 5 Increasing Risk Factors": ["", ".join([f for f in feat if f != "])",
    ↪for feat in top_increase_features],
    "Top 5 Decreasing Risk Factors": ["", ".join([f for f in feat if f != "])",
    ↪for feat in top_decrease_features]
})

# Set pandas display options for left alignment
pd.set_option('display.max_colwidth', 50)
pd.set_option('display.width', 120)
pd.set_option('display.colheader_justify', 'left')

# Display FIRST 10 patients
print("FIRST 10 PATIENTS")
for i in range(min(10, len(results_table))):
    row = results_table.iloc[i]
    print(f"Patient: {row['Patient']:<5} | Diagnosis: ",
    ↪{row['Diagnosis(1-Anxiety, 0-No Anxiety)']:<13} | "
    f"Prob: {row['Predicted Prob (%)']:<6}% | Risk: {row['Risk Category']:<12}")
    print(f"  ↑ Increasing: {row['Top 5 Increasing Risk Factors']}")
    print(f"  ↓ Decreasing: {row['Top 5 Decreasing Risk Factors']}")
    print("-" * 110)

# Display LAST 10 patients
print("LAST 10 PATIENTS")
start_idx = max(0, len(results_table) - 10)

```

```

for i in range(start_idx, len(results_table)):
    row = results_table.iloc[i]
    print(f"Patient: {row['Patient']:<5} | Diagnosis:␣
↪{row['Diagnosis(1-Anxiety, 0-No Anxiety)']:<13} | "
        f"Prob: {row['Predicted Prob (%)']:<6}% | Risk: {row['Risk Category']:
↪<12}")
    print(f"    ↑ Increasing: {row['Top 5 Increasing Risk Factors']}")
    print(f"    ↓ Decreasing: {row['Top 5 Decreasing Risk Factors']}")
    print("-" * 110)

# Save full results table to CSV
results_table.to_csv("predictions_with_explanation.csv", index=False)
print("\nFull results table saved to 'predictions_with_explanations.csv'")

# Create summary statistics for risk categories
print("RISK CATEGORY DISTRIBUTION")
risk_summary = results_table.groupby('Risk Category').agg({
    'Patient': 'count',
}).rename(columns={'Patient': 'Count'})
risk_summary['Percentage'] = (risk_summary['Count'] / len(results_table) * 100).
↪round(2)
print(risk_summary.to_string())

# Save model
joblib.dump(xgb_model, 'predictive_model.pkl', compress=3)
print("\nModel saved as predictive_model.pkl")


```



/workspaces/DECLINE/.venv/lib/python3.12/site-packages/tqdm/auto.py:21:
TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

Predicted percentages (first 20 patients):

Patient	Diagnosis(1-Anxiety, 0-No Anxiety)	Logistic Regression (%)	Random Forest (%)	XGBoost (%)
1		1	79.31	
78.96	73.27			
2		0	0.30	
1.16	0.62			
3		0	0.30	
4.07	0.57			
4		0	5.63	
13.09	9.71			
5		0	0.00	
0.18	0.00			
6		0	58.51	
65.57	63.97			

7	0	21.99
28.99	28.32	
8	0	17.32
18.73	15.44	
9	0	1.68
2.32	2.35	
10	0	17.63
19.26	11.93	
11	0	1.02
3.14	5.94	
12	0	0.19
2.35	0.29	
13	0	41.44
41.93	30.55	
14	0	35.12
35.41	27.32	
15	0	11.78
8.19	13.90	
16	0	0.76
0.00	1.66	
17	0	7.78
1.16	10.17	
18	0	26.54
24.08	22.17	
19	0	12.32
15.77	21.92	
20	1	98.54
100.00	99.71	






Top 15 features for Logistic Regression:


Stress Level (1-10): 1.8771 (increases risk)
 Sleep Hours: -0.7585 (decreases risk)
 Therapy Sessions (per month): 0.5316 (increases risk)
 Caffeine Intake (mg/day): 0.4717 (increases risk)
 Occupation_Student: -0.1800 (decreases risk)
 Occupation_Nurse: 0.1441 (increases risk)
 Occupation_Doctor: 0.1422 (increases risk)
 Occupation_Other: 0.1068 (increases risk)
 Medication: 0.0897 (increases risk)
 Recent Major Life Event: 0.0809 (increases risk)
 Occupation_Lawyer: 0.0807 (increases risk)

Occupation_Athlete: -0.0805 (decreases risk)
Alcohol Consumption (drinks/week): 0.0739 (increases risk)
Gender_Other: 0.0732 (increases risk)
Age: -0.0731 (decreases risk)




Top 15 features for Random Forest:

Stress Level (1-10): 0.3114
Sleep Hours: 0.1086
Caffeine Intake (mg/day): 0.0850
Physical Activity (hrs/week): 0.0589
Therapy Sessions (per month): 0.0552
Heart Rate (bpm): 0.0536
Age: 0.0509
Breathing Rate (breaths/min): 0.0438
Alcohol Consumption (drinks/week): 0.0428
Diet Quality (1-10): 0.0404
Sweating Level (1-5): 0.0271
Recent Major Life Event: 0.0100
Medication: 0.0096
Dizziness: 0.0095
Gender_Other: 0.0089



Top 15 features for XGBoost:

Stress Level (1-10): 0.4402
Sleep Hours: 0.0664
Therapy Sessions (per month): 0.0482
Caffeine Intake (mg/day): 0.0319
Gender_Male: 0.0246
Occupation_Athlete: 0.0218
Occupation_Chef: 0.0201
Occupation_Scientist: 0.0191
Sweating Level (1-5): 0.0180
Alcohol Consumption (drinks/week): 0.0176
Occupation_Musician: 0.0176
Diet Quality (1-10): 0.0166
Occupation_Lawyer: 0.0164
Age: 0.0162
Physical Activity (hrs/week): 0.0162



Logistic Regression METRICS (threshold=0.284)

Confusion Matrix:

True Negatives: 6103
 False Positives: 1515
 False Negatives: 573
 True Positives: 2809

Key Performance Metrics:

- SENSITIVITY (Recall/True Positive Rate): 0.8306 (83.06%)
- SPECIFICITY (True Negative Rate): 0.8011 (80.11%)
- F1-SCORE: 0.7290

Additional Metrics:

Accuracy: 0.8102 (81.02%)
 Precision: 0.6496 (64.96%)
 ROC AUC: 0.9047
 Brier Score: 0.1114

Classification Report:

	precision	recall	f1-score	support
0	0.9142	0.8011	0.8539	7618
1	0.6496	0.8306	0.7290	3382

accuracy			0.8102	11000
macro avg	0.7819	0.8159	0.7915	11000
weighted avg	0.8328	0.8102	0.8155	11000

Random Forest METRICS (threshold=0.293)

Confusion Matrix:

True Negatives: 5923
False Positives: 1695
False Negatives: 510
True Positives: 2872

Key Performance Metrics:

- SENSITIVITY (Recall/True Positive Rate): 0.8492 (84.92%)
- SPECIFICITY (True Negative Rate): 0.7775 (77.75%)
- F1-SCORE: 0.7226

Additional Metrics:

Accuracy: 0.7995 (79.95%)
Precision: 0.6289 (62.89%)
ROC AUC: 0.8993
Brier Score: 0.1136

Classification Report:

	precision	recall	f1-score	support
0	0.9207	0.7775	0.8431	7618
1	0.6289	0.8492	0.7226	3382
accuracy			0.7995	11000
macro avg	0.7748	0.8134	0.7828	11000
weighted avg	0.8310	0.7995	0.8060	11000

XGBoost METRICS (threshold=0.267)

Confusion Matrix:

True Negatives: 5787
False Positives: 1831
False Negatives: 473
True Positives: 2909

Key Performance Metrics:

- SENSITIVITY (Recall/True Positive Rate): 0.8601 (86.01%)
- SPECIFICITY (True Negative Rate): 0.7596 (75.96%)
- F1-SCORE: 0.7163

Additional Metrics:

Accuracy: 0.7905 (79.05%)
Precision: 0.6137 (61.37%)
ROC AUC: 0.9000


Brier Score: 0.1136



Classification Report:

	precision	recall	f1-score	support
0	0.9244	0.7596	0.8340	7618
1	0.6137	0.8601	0.7163	3382
accuracy			0.7905	11000
macro avg	0.7691	0.8099	0.7752	11000
weighted avg	0.8289	0.7905	0.7978	11000

MODEL COMPARISON SUMMARY


	Model	Threshold	Sensitivity (%)	Specificity (%)	F1-Score
Precision (%)	Accuracy (%)	ROC-AUC			
Logistic Regression	0.28		83.06	80.11	0.73
64.96	81.02	0.9			
Random Forest	0.29		84.92	77.75	0.72
62.89	79.95	0.9			
XGBoost	0.27		86.01	75.96	0.72
61.37	79.05	0.9			







Generating SHAP Waterfall Plots for First 3 Patients...

=====







RESULTS TABLE

FIRST 10 PATIENTS

Patient: 1 | Diagnosis: Diagnosed | Prob: 73.27 % | Risk: High Risk

↑ Increasing: Stress Level (1-10) (+2.160), Age (+0.130), Heart Rate (bpm) (+0.059), Recent Major Life Event (+0.057), Breathing Rate (breaths/min) (+0.055)

↓ Decreasing: Therapy Sessions (per month) (-0.399), Caffeine Intake (mg/day) (-0.233), Gender_Other (-0.022), Dizziness (-0.014), Smoking (-0.007)

Patient: 2 | Diagnosis: Not Diagnosed | Prob: 0.62 % | Risk: Low Risk

↑ Increasing: Physical Activity (hrs/week) (+0.308), Family History of Anxiety (+0.085), Sleep Hours (+0.047), Occupation_Nurse (+0.039), Occupation_Student (+0.018)

↓ Decreasing: Stress Level (1-10) (-3.806), Therapy Sessions (per month) (-0.542), Caffeine Intake (mg/day) (-0.468), Heart Rate (bpm) (-0.094), Gender_Other (-0.084)

Patient: 3 | Diagnosis: Not Diagnosed | Prob: 0.57 % | Risk: Low Risk
↑ Increasing: Sleep Hours (+0.580), Heart Rate (bpm) (+0.097), Family History
of Anxiety (+0.089), Recent Major Life Event (+0.077), Gender_Other (+0.051)
↓ Decreasing: Stress Level (1-10) (-3.843), Therapy Sessions (per month)
(-0.943), Caffeine Intake (mg/day) (-0.625), Physical Activity (hrs/week)
(-0.138), Diet Quality (1-10) (-0.116)

Patient: 4 | Diagnosis: Not Diagnosed | Prob: 9.71 % | Risk: Low Risk
↑ Increasing: Sleep Hours (+0.439), Occupation_Student (+0.020),
Occupation_Chef (+0.010), Caffeine Intake (mg/day) (+0.005), Diet Quality (1-10)
(+0.004)
↓ Decreasing: Stress Level (1-10) (-1.456), Therapy Sessions (per month)
(-0.811), Recent Major Life Event (-0.116), Physical Activity (hrs/week)
(-0.085), Alcohol Consumption (drinks/week) (-0.059)

Patient: 5 | Diagnosis: Not Diagnosed | Prob: 0.0 % | Risk: Low Risk
↑ Increasing: Occupation_Other (+0.113), Medication (+0.054), Gender_Other
(+0.040), Breathing Rate (breaths/min) (+0.038), Dizziness (+0.025)
↓ Decreasing: Stress Level (1-10) (-4.251), Therapy Sessions (per month)
(-0.871), Sleep Hours (-0.661), Caffeine Intake (mg/day) (-0.278), Alcohol
Consumption (drinks/week) (-0.123)

Patient: 6 | Diagnosis: Not Diagnosed | Prob: 63.97 % | Risk: Moderate Risk
↑ Increasing: Stress Level (1-10) (+1.534), Physical Activity (hrs/week)
(+0.119), Recent Major Life Event (+0.067), Occupation_Other (+0.064),
Medication (+0.063)
↓ Decreasing: Therapy Sessions (per month) (-0.618), Age (-0.139), Alcohol
Consumption (drinks/week) (-0.069), Sleep Hours (-0.062), Occupation_Doctor
(-0.009)

Patient: 7 | Diagnosis: Not Diagnosed | Prob: 28.32 % | Risk: Low Risk
↑ Increasing: Stress Level (1-10) (+1.042), Heart Rate (bpm) (+0.230), Age
(+0.127), Physical Activity (hrs/week) (+0.083), Recent Major Life Event
(+0.066)
↓ Decreasing: Therapy Sessions (per month) (-0.756), Caffeine Intake (mg/day)
(-0.680), Sleep Hours (-0.474), Medication (-0.028), Gender_Other (-0.020)

Patient: 8 | Diagnosis: Not Diagnosed | Prob: 15.44 % | Risk: Low Risk
↑ Increasing: Caffeine Intake (mg/day) (+0.594), Physical Activity (hrs/week)
(+0.156), Alcohol Consumption (drinks/week) (+0.150), Sleep Hours (+0.076),
Breathing Rate (breaths/min) (+0.035)

↓ Decreasing: Stress Level (1-10) (-0.973), Therapy Sessions (per month) (-0.698), Age (-0.117), Heart Rate (bpm) (-0.081), Medication (-0.067)

Patient: 9 | Diagnosis: Not Diagnosed | Prob: 2.35 % | Risk: Low Risk
↑ Increasing: Caffeine Intake (mg/day) (+0.691), Sleep Hours (+0.130), Heart Rate (bpm) (+0.075), Occupation_Lawyer (+0.073), Recent Major Life Event (+0.050)

↓ Decreasing: Stress Level (1-10) (-2.967), Therapy Sessions (per month) (-0.597), Alcohol Consumption (drinks/week) (-0.131), Age (-0.109), Physical Activity (hrs/week) (-0.100)

Patient: 10 | Diagnosis: Not Diagnosed | Prob: 11.93 % | Risk: Low Risk
↑ Increasing: Sleep Hours (+0.696), Recent Major Life Event (+0.095), Alcohol Consumption (drinks/week) (+0.085), Medication (+0.052), Sweating Level (1-5) (+0.048)

↓ Decreasing: Stress Level (1-10) (-1.152), Therapy Sessions (per month) (-0.933), Caffeine Intake (mg/day) (-0.514), Age (-0.145), Heart Rate (bpm) (-0.098)

LAST 10 PATIENTS

Patient: 10991 | Diagnosis: Not Diagnosed | Prob: 0.0 % | Risk: Low Risk
↑ Increasing: Gender_Other (+0.058), Medication (+0.053), Heart Rate (bpm) (+0.037), Dizziness (+0.022), Occupation_Student (+0.017)

↓ Decreasing: Stress Level (1-10) (-3.556), Therapy Sessions (per month) (-0.866), Sleep Hours (-0.630), Caffeine Intake (mg/day) (-0.233), Physical Activity (hrs/week) (-0.118)

Patient: 10992 | Diagnosis: Not Diagnosed | Prob: 5.09 % | Risk: Low Risk
↑ Increasing: Caffeine Intake (mg/day) (+0.809), Alcohol Consumption (drinks/week) (+0.233), Occupation_Doctor (+0.108), Age (+0.072), Gender_Other (+0.036)

↓ Decreasing: Stress Level (1-10) (-2.946), Therapy Sessions (per month) (-0.508), Heart Rate (bpm) (-0.165), Medication (-0.086), Sleep Hours (-0.060)

Patient: 10993 | Diagnosis: Not Diagnosed | Prob: 31.58 % | Risk: Moderate Risk
↑ Increasing: Stress Level (1-10) (+0.366), Recent Major Life Event (+0.099), Alcohol Consumption (drinks/week) (+0.076), Medication (+0.052), Breathing Rate (breaths/min) (+0.042)

↓ Decreasing: Therapy Sessions (per month) (-0.467), Caffeine Intake (mg/day) (-0.108), Age (-0.082), Physical Activity (hrs/week) (-0.076), Sleep Hours (-0.058)

Patient: 10994 | Diagnosis: Not Diagnosed | Prob: 0.19 % | Risk: Low Risk
↑ Increasing: Alcohol Consumption (drinks/week) (+0.108), Medication (+0.051), Gender_Other (+0.049), Breathing Rate (breaths/min) (+0.032), Dizziness (+0.023)
↓ Decreasing: Stress Level (1-10) (-3.635), Sleep Hours (-0.794), Caffeine Intake (mg/day) (-0.193), Recent Major Life Event (-0.101), Physical Activity (hrs/week) (-0.074)

Patient: 10995 | Diagnosis: Not Diagnosed | Prob: 13.8 % | Risk: Low Risk
↑ Increasing: Therapy Sessions (per month) (+0.533), Diet Quality (1-10) (+0.044), Sweating Level (1-5) (+0.039), Heart Rate (bpm) (+0.025), Medication (+0.023)
↓ Decreasing: Stress Level (1-10) (-1.409), Caffeine Intake (mg/day) (-0.486), Recent Major Life Event (-0.093), Breathing Rate (breaths/min) (-0.092), Age (-0.077)

Patient: 10996 | Diagnosis: Diagnosed | Prob: 55.57 % | Risk: Moderate Risk
↑ Increasing: Stress Level (1-10) (+0.886), Caffeine Intake (mg/day) (+0.578), Age (+0.293), Heart Rate (bpm) (+0.100), Medication (+0.099)
↓ Decreasing: Therapy Sessions (per month) (-0.837), Physical Activity (hrs/week) (-0.122), Sweating Level (1-5) (-0.082), Recent Major Life Event (-0.074), Diet Quality (1-10) (-0.044)

Patient: 10997 | Diagnosis: Not Diagnosed | Prob: 17.66 % | Risk: Low Risk
↑ Increasing: Stress Level (1-10) (+0.325), Alcohol Consumption (drinks/week) (+0.068), Sweating Level (1-5) (+0.063), Occupation_Teacher (+0.018), Gender_Other (+0.015)
↓ Decreasing: Caffeine Intake (mg/day) (-0.551), Therapy Sessions (per month) (-0.470), Heart Rate (bpm) (-0.123), Age (-0.110), Recent Major Life Event (-0.108)

Patient: 10998 | Diagnosis: Not Diagnosed | Prob: 46.28 % | Risk: Moderate Risk
↑ Increasing: Stress Level (1-10) (+0.924), Physical Activity (hrs/week) (+0.271), Age (+0.100), Medication (+0.097), Occupation_Nurse (+0.067)
↓ Decreasing: Therapy Sessions (per month) (-0.516), Caffeine Intake (mg/day) (-0.376), Sleep Hours (-0.134), Heart Rate (bpm) (-0.073), Sweating Level (1-5) (-0.032)

Patient: 10999 | Diagnosis: Not Diagnosed | Prob: 9.57 % | Risk: Low Risk
↑ Increasing: Sleep Hours (+0.560), Recent Major Life Event (+0.079), Sweating Level (1-5) (+0.036), Physical Activity (hrs/week) (+0.036), Gender_Other (+0.032)
↓ Decreasing: Stress Level (1-10) (-1.551), Therapy Sessions (per month) (-0.923), Caffeine Intake (mg/day) (-0.154), Age (-0.144), Alcohol Consumption

(drinks/week) (-0.066)

Patient: 11000 | Diagnosis: Not Diagnosed | Prob: 0.42 % | Risk: Low Risk
 ↑ Increasing: Breathing Rate (breaths/min) (+0.157), Recent Major Life Event
 (+0.093), Sleep Hours (+0.070), Alcohol Consumption (drinks/week) (+0.059),
 Sweating Level (1-5) (+0.045)
 ↓ Decreasing: Stress Level (1-10) (-3.913), Therapy Sessions (per month)
 (-0.546), Caffeine Intake (mg/day) (-0.284), Gender_Other (-0.117), Heart Rate
 (bpm) (-0.111)

Full results table saved to 'predictions_with_explanations.csv'

RISK CATEGORY DISTRIBUTION

	Count	Percentage
Risk Category		
High Risk	1386	12.60
Low Risk	6614	60.13
Moderate Risk	3000	27.27

Model saved as predictive_model.pkl